

Evaluation of Custom Embedding Model g_1

Overview

We aim to fine-tune a CALE embedding model to better associate cryptic crossword definitions with their correct answers. Our first custom model, g_1 , was trained using triplet margin loss ($\alpha = 1.0$, $p = 2, 3$ epochs, $\text{lr} = 2 \times 10^{-5}$, batch size 32) on 69,921 training triplets (τ_1), each consisting of:

Anchor	$f_{\text{clue}}(\text{def})$	the definition word, embedded in its cryptic clue context
Positive	$f_{\text{wndef}}(\text{ans})$	the correct answer, embedded with its WordNet definition
Negative	$f_{\text{wndef}}(\text{dist})$	a distractor word, embedded with its WordNet definition

The training objective pushes the anchor closer to the positive and farther from the negative. We evaluated g_1 on 46,506 held-out validation triplets using both the training metric (L2 distance) and the research metric (cosine similarity).

Data Leakage from Seen Vocabulary

While every anchor (clue-contextualized definition) in the validation set is unseen, positives and negatives may appear in both the training and validation splits—an unavoidable feature of language reusing words. Where appropriate, we stratify analyses by *seen* and *unseen* vocabulary, and use the f_{wndef} phrase construction strategy for decontextualized meanings.

Metric Mismatch: A Non-Issue

We trained g_1 with L2 distance but evaluated with cosine similarity. L2 and cosine agree on every analysis in this evaluation—triplet accuracy, margins, context effects, and stratified comparisons all produce the same qualitative findings under both metrics.¹ The choice of training metric did not distort any of our conclusions, so below we report metrics only in cosine similarity.

Task Performance and Overfitting

The model clearly learned the triplet discrimination task: g_1 achieves 90.0% triplet accuracy on held-out validation triplets, compared to g_{stock} ’s 38.8%.² The *margin*—the difference between the anchor-positive similarity and the anchor-negative similarity for each triplet—tells us how well-separated the correct answer is from the distractor. g_1 ’s validation mean margin is 0.125, compared to g_{stock} ’s -0.054 . The negative g_{stock} margins confirm that distractors are typically more similar to the anchor than answers are under the pretrained model.

¹This is expected when embedding norms are tightly concentrated: L2 distance depends on both direction and magnitude, while cosine depends only on direction, so the two converge when magnitudes are similar. Under g_1 , the **wndef** norm distribution has a coefficient of variation of $\sim 3\%$ (std 0.86 on a mean of 27.9), and **f_clue** is even tighter at $\sim 2.6\%$.

²The fact that g_{stock} scores below chance (50%) is itself informative: the τ_1 distractors are systematically more similar to the anchor than the true answers are under pretrained g_{stock} . The training asked g_1 to invert this ordering.




Figure 1: Training dynamics across 3 epochs. Left: training loss (green) diverges from validation loss (red) after epoch 2. Center: validation accuracy plateaus near 90%. Right: validation margins continue growing even as loss rises.

Training dynamics show classic overfitting: validation loss improved from epoch 1 to epoch 2 ($0.305 \rightarrow 0.252$) but rose at epoch 3 (0.264), while training loss fell to 0.014 —a $19\times$ train/val ratio. Notably, validation accuracy barely changed from epoch 2 to 3 ($89.7\% \rightarrow 90.0\%$) while validation margins grew substantially ($0.088 \rightarrow 0.125$). The model polarized the validation set: triplets it already classified correctly were pushed further apart, while the $\sim 10\%$ it got wrong incurred larger penalties—which is why loss rose even as margins grew.

The gap between training accuracy ($\geq 98.6\%$ inferred from epoch 3 training loss) and validation accuracy (90.0%) confirms that g_1 performs better on data it trained on. But this gap is modest, and stratifying the validation triplets by whether the positive and negative words appeared in training tells us it didn’t produce word-level memorization:




Figure 2: Cosine triplet accuracy by vocabulary exposure stratum, with 95% confidence intervals. Both g_{stock} and g_1 perform uniformly across strata. The original model is below chance (dashed line) in all three.

g_1 's performance is uniform across seen and unseen words; the overfitting likely means we trained one epoch past optimal, but it didn't produce data leakage at the word level.

Embedding Space: Compressed and Reorganized

Fine-tuning contracted the embedding space substantially without collapsing it dimensionally:

		g_{stock}	g_1	Shift
wndef	Mean pairwise cosine	0.412	0.561	+0.149
	Total variance	27.0M	18.6M	-31%
	Effective dimensionality	41.1	45.4	+4.4
	RSA ρ	—	0.135	—
f_clue	Mean pairwise cosine	0.317	0.452	+0.135
	Total variance	30.7M	22.1M	-28%
	Effective dimensionality	52.3	103.2	+50.8
	RSA ρ	—	0.067	—

The RSA measures whether words that were similar under g_{stock} remain similar under g_1 . At $\rho = 0.14$ for **wndef** and $\rho = 0.07$ for **f_clue**, g_1 didn't refine the pretrained similarity structure—it essentially discarded it and built something new. The pairwise cosine increase (+0.15) means unrelated words became more similar to each other, undermining discriminability.

Total variance decreased while effective dimensionality increased—especially for **f_clue**, where it nearly doubled ($52 \rightarrow 103$). This means g_1 compressed the dominant directions of variation while spreading what remained more evenly across dimensions. The space got smaller but flatter: it did not collapse, but it lost some of the structure of the original CALE model.

This compression is not localized to words the model trained on. Stratifying the vocabulary by training exposure shows that seen and unseen words were affected equally:

	Mean pairwise cosine	
	g_{stock}	g_1
Seen ($n = 36,082$)	0.408	0.565
Unseen ($n = 12,794$)	0.431	0.557

The model applied a global transformation to the entire embedding space, including regions it never directly optimized—a structural change to how the model maps text to embedding space, not a local side effect of training on specific words.

No Evidence of Cryptic Crossword Specific Learning

The central question is whether g_1 learned to extract useful information from cryptic clue context. $T = 0$ measures decontextualized definition-answer similarity (definition appears alone); $T = 1$ measures clue-contextualized definition-answer similarity (definition appears inside its cryptic clue). The ATE is $\overline{T=1} - \overline{T=0}$; a negative value means clue context pushed the definition farther from the answer.




Figure 3: Cosine similarity distributions for $T=0$ (definition alone) and $T=1$ (definition in clue context). Under g_{stock} (left), the distributions overlap substantially. Under g_1 (right), both distributions tightened and shifted right, but $T=0$ shifted much more than $T=1$, widening the gap between them.

	g_{stock}	g_1
$T=0$ (definition alone)	0.576	0.715
$T=1$ (definition in clue)	0.513	0.591
ATE ($T=1 - T=0$)	-0.063	-0.124
% pairs where context hurts	71.3%	93.4%

Under g_1 , decontextualized definition-answer pairs became closer together ($\Delta T = 0$: +0.139), as did the clue-contextualized definition-answer pairs but to a much lesser extent ($\Delta T = 1$: +0.078). The ATE doubled in magnitude, and the fraction of pairs where context is misleading rose from 71% to 93%.

The custom model did not learn to exploit cryptic clue structure. Under g_1 , clue context became even more misleading.

Summary

Customizing the model fundamentally reorganized the embedding space, made random words more similar to each other, and amplified rather than reduced the misdirection from clue context. The next set of investigations will test which methodological issues—distractor design, phrase format, noun bias, sense selection—contributed to this behavior.